

SHINOJ PHILIP JOHN

Melbourne, VIC | +61 456 995 237 | shinojphilipjohn@gmail.com | [linkedin.com/in/shinoj-philip-john](https://www.linkedin.com/in/shinoj-philip-john) | www.shinojphilipjohn.com

PROFESSIONAL SUMMARY

Data engineer with hands-on experience designing, building, and operating production-grade ETL/ELT pipelines, cloud data platforms, and modern data stack solutions across AWS and Azure environments. Proficient in Python, SQL, dbt Core, Snowflake, and Databricks, with a strong foundation in Medallion Architecture, data quality testing, CI/CD, and pipeline orchestration. Holds a Databricks Fundamentals certification, Microsoft Azure Data Fundamentals (DP-900) in progress, and a Master of Analytics from RMIT University — committed to delivering scalable, reliable data infrastructure that powers business decisions.

SKILLS

Programming & Tools: Python, SQL, dbt Core, Azure Data Factory, GitHub Actions (CI/CD), SQLAlchemy, pandas, NumPy, Scikit-learn, Matplotlib, Jupyter, Git, R (basic)

Cloud & Data: Azure (Blob Storage, SQL Database, Databricks, Entra ID), AWS (S3, EMR/Apache Spark, Redshift, Lambda), Snowflake, Medallion Architecture, ETL/ELT Development, Data Pipelines, Data Modelling, Data Quality Testing (dbt, SQL Validation), Data Warehousing, Incremental Loading (CDC), IBM Workload Scheduler, YAML/JSON Configuration

Data Science & ML: Supervised ML (Logistic Regression, Random Forest), EDA, Feature Engineering, Model Evaluation, Predictive Analytics, Statistical Modelling, Experiment Design

Visualization & Storytelling: Power BI, Tableau, Data Narratives, Communication for Non-technical Stakeholders

Soft Skills: Analytical thinking, structured problem-solving, collaboration, consulting communication

WORK EXPERIENCE

Data Engineer (Internship), IBISWorld | 2025 - 2026

- Delivered Snowflake ELT workflows using AWS S3 external tables, improving ingestion efficiency and reducing unnecessary storage duplication by **~20–25%** across recurring loads.
- Parsed and flattened complex multi-level JSON using **LATERAL FLATTEN**, increasing usable analytical fields by **~35%** and cutting manual preprocessing time by **~30%** through a structured multi-layer data flow (staging → parsing → transformation).
- Enhanced pipeline traceability with **structured logging** and clear transformation checkpoints, reducing debugging time by **~25%** and supporting faster onboarding through high-level architecture documentation.

Data Engineer, nbm Australia (via Infosys) | 2021–2023

- Built and maintained AWS S3 → EMR → Redshift ETL pipelines supporting financial and reporting teams, achieving **~99% successful daily run stability** across batch workloads.
- Optimised SQL transformations and Redshift table structures (sort/distribution keys), improving end-to-end pipeline runtime by **~25–30%** on recurring data loads.
- Developed and managed **50+ SQL scripts** including delta-load logic, DDLs, and control-table updates in MySQL, ensuring alignment with E-R designs and JIRA user stories.
- Reduced manual intervention during executions by **~40%** through parameterising ETL behaviour using YAML/JSON files and automating jobstreams in IBM Workload Scheduler.
- Partnered with reporting and finance stakeholders to translate business requirements into SQL transformations and consumable datasets.
- Enhanced data quality and reduced SIT→UAT defect leakage by **~30%** by validating missing data, identifying upstream design gaps, and coordinating fixes with analysts and operations teams.

PROJECTS

Brazilian E-Commerce ELT Pipeline (Azure, Python, dbt Core, ADF)

- Engineered a production-grade ELT pipeline ingesting the Olist Brazilian E-Commerce dataset (**~100K orders, 9 relational tables**) from Azure Blob Storage into Azure SQL Database using Python, azure-storage-blob SDK, and

SQLAlchemy with Microsoft Entra DefaultAzureCredential authentication — **zero hardcoded credentials** and file-level error handling ensuring individual failures are caught without stopping the pipeline.

- Transformed raw data through a **Medallion Architecture (raw → staging → mart)** using dbt Core with dbt-sqlserver adapter, implementing staging models, **Kimball dimensional mart tables**, incremental materialisation, schema-level data quality tests (unique, not_null, relationships), composite key validation via dbt_utils, and dev/prod environment separation via profiles.yml.
- Orchestrated daily ingestion via **Azure Data Factory** with **GitHub Actions CI/CD** executing dbt run and dbt test on every push, pull request, and daily scheduled cron — branch protection rules enforce that only tested code reaches production, ensuring data quality on every deployment at **AU\$1.01 infrastructure cost** on serverless Azure SQL.

Customer Churn Prediction (Python, Scikit-learn)

- Cleaned and preprocessed **7,000+** customer records, handling missing values, encoding categorical fields, and scaling numerical features to prepare data for modelling.
- Conducted EDA to uncover trends in tenure, contract types, and billing patterns, enabling more informed feature selection.
- Improved model accuracy from a baseline of **~70%** to **~82%** by applying feature engineering, hyperparameter tuning, and cross-validation.

EDUCATION

Master of Analytics (STEM), RMIT University | 2023–2025

B.Tech– Electrical & Electronics Engineering, APJ Abdul Kalam Technological University | 2016–2020

CERTIFICATIONS

Microsoft Azure Data Fundamentals (DP-900) – Currently Completing

Academy Accreditation – Databricks Fundamentals | Databricks